

NAMAN MANOCHA

GENERATIVE AI ENGINEER — LLMs, RAG & Prompt Engineering

Kolkata, India | 8530356515 | namanmanocha42248@gmail.com | [Linkedin](#) | [GitHub](#)

SUMMARY

GenAI engineer focused on grounded LLM applications: retrieval, citation, and abstention over confident hallucination. Solo-developed FundersAI, a citation-grounded seven-stage LangGraph RAG workflow, and CareFlow Intelligence, which pairs agentic RAG with a RAFT fine-tuning workflow. Core technologies: LangGraph, OpenAI/Groq, pgvector, Unsloth, and prompt engineering.

TECHNICAL SKILLS

LLM & GenAI: LangGraph, Google ADK, RAG, Prompt Engineering, Hybrid Search, Reranking, pgvector, ChromaDB, Hugging Face, Ollama, Langfuse

Model Development: MLflow, scikit-learn, K-Means Clustering, Logistic Regression, PyTorch, BERT Fine-Tuning

Backend & Cloud: Python, FastAPI, REST APIs, SQL, PostgreSQL, Supabase, Docker, Kubernetes, GitHub Actions, Next.js, TypeScript

EXPERIENCE

Technology Manager — PayGain Multiservices Pvt. Ltd.

Feb 2025 – Present

Kolkata, India

- Scoping and leading development of a company-wide mobile app as sole engineer; also designed a lead-capture automation (Make) triggered on Google Form submissions, eliminating manual lead logging.

INDEPENDENT PROJECTS

FundersAI — AI-Powered Research Platform for Indian Stocks & Mutual Funds

fundersai.co.in

Python, FastAPI, Next.js, Supabase/pgvector, LangGraph, Groq, OpenAI | ~143K source LOC, 856 files, 442 commits, 178 test files — OpenAI Build Week | Live: www.fundersai.co.in | Repo: github.com/Naman6019/FundersAI

- Citation-Grounded RAG: built a seven-stage LangGraph workflow with retrieval grading, a query-rewrite loop, lexical reranking, citation-aware synthesis, and strict abstention when evidence is insufficient.
- Multi-Agent Orchestration:** designed a custom multi-agent pipeline (Planner → Quant → Research → Synthesis agents) with shared state, where a Groq-hosted LLM router agent classifies query intent and directs execution across stock, mutual-fund, comparison, and market-news paths.
- Prompt Engineering:** iterated few-shot/CoT prompts to control synthesis and abstention behavior across the pipeline.

CareFlow Intelligence — Agentic RAG & Fine-Tuning System

Python, FastAPI, Next.js, Supabase/pgvector, ChromaDB, Unsloth, Ollama | ~7,300 LOC, non-clinical, synthetic-data-only

- Agentic Citation-Grounded RAG: built a bounded-tool document agent with query reformulation and mandatory citation validation against PostgreSQL full-text indexes, backed by automated OpenFDA, PubMed, guideline, and media ingestion.
- RAFT Fine-Tuning: generated training data with reference evidence, hard distractors, and abstention cases; fine-tuned Qwen2.5-1.5B via 4-bit QLoRA (rank 16, alpha 16), then quantized it to GGUF and served it locally via Ollama.

Other Projects

- TalentOS (All Things Agentic Hackathon, Taskmaster Track): dual-pipeline opportunity platform using Google ADK, LangGraph, and Gemini on Vertex AI to generate human-reviewable resumes and pitches after deterministic pre-filtering (89.5% of ~2,900 postings/run eliminated before any LLM call), validated by a 373-test hermetic offline suite (~11s). FundersAI Reports: decoupled LangGraph report-synthesis service hosted on K3s/EC2. SQuAD QA Benchmarking: fine-tuned BERT, BiDAF, and DistilBERT; published as a research paper. GitResume: self-built evidence-backed resume CLI (PyPI: git-resume-agent, ~4K LOC, 28 commits) that validates managed claims before compiling Word/PDF output.

EDUCATION & ACHIEVEMENTS

B.Tech, Computer Science Engineering (AI & ML Specialization) — VIT Amaravati

June 2021 – Sept 2025

CGPA: 8.2/10 | Published research paper on QA/NLP | Certifications: Azure Administrator Associate, AWS Certified Cloud Practitioner