

NAMAN MANOCHA

AI ENGINEER • GENAI • AGENTIC SYSTEMS — LLMs, RAG & Applied ML
Kolkata, India | 8530356515 | namanmanocha42248@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

AI/ML engineer with experience across classical ML, multi-agent orchestration, and grounded RAG. Independently developed FundersAI, TalentOS, and CareFlow using Python, FastAPI, Next.js, Docker, Kubernetes, and pgvector.

TECHNICAL SKILLS

ML & AI: LangGraph, Google ADK, RAG, Hybrid Search, pgvector, QLoRA, Prompt Engineering, scikit-learn, MLflow, Langfuse

Backend & Cloud: Python, FastAPI, REST APIs, SQL, PostgreSQL, Supabase, Docker, Kubernetes, Google Cloud Run, GitHub Actions

Web & Data: Next.js, TypeScript, ETL Pipelines, Data Validation

EXPERIENCE

Artificial Intelligence Intern — Internmo (Remote)

Sept 2026 – Present

Client-style ML projects: Vantara (churn/CLV, SHAP, FastAPI, Docker), FlowCast (traffic and accident-risk forecasting), Northgate (FinBERT, max-Sharpe).

Technology Manager — PayGain Multiservices Pvt. Ltd.

Feb 2025 – Sept 2026

Kolkata, India

Scoped and designed a company-wide mobile app (design stage), covering client/employee onboarding, document access, product activation, and payments — mobile counterpart to PayGain core platform.

Designed and deployed lead-capture automation (Make); drove 30+ client conversions via website and digital discovery flows; simplified complex financial products for 300+ seminar attendees via 20+ technical presentations.

INDEPENDENT PROJECTS

FundersAI — AI-Powered Research Platform for Indian Stocks & Mutual Funds

fundersai.co.in

Python, FastAPI, Next.js, Supabase/pgvector, LangGraph, Groq, OpenAI, Cloudflare R2, Kubernetes (K3s)

Solo-built, Apr–Aug 2026 — ~144K source LOC across 858 files, 443 commits, 179 test files — OpenAI Build Week | Live:

www.fundersai.co.in | Repo: github.com/Naman6019/FundersAI

Citation-Grounded RAG & Multi-Agent Routing: built a seven-stage LangGraph workflow with evidence-aware synthesis, abstention, lexical reranking, and routing across stock, fund, and news paths.

Explainable ML & Microservice Architecture: engineered fund-similarity clustering (k-means, 13 features) and logistic regression triage; extracted report synthesis into a decoupled FastAPI/LangGraph microservice deployed on a 2-replica K3s cluster on AWS EC2.

Resilient Data Engineering: automated official-disclosure collection and normalization with readiness checks and targeted normalizer repairs.

Cloud Storage & Billing: built an S3-compatible storage layer (Cloudflare R2, boto3/AWS SigV4) for raw documents and cold archival; shipped Razorpay subscription billing (webhook-verified, Free/Pro/Ultra tiers) with Supabase RLS enforcement.

TalentOS — Autonomous Opportunity Intelligence Platform (All Things Agentic Hackathon)

Python, Google ADK, LangGraph, Next.js, Firestore, Google Cloud Run | Solo-built, Aug 2026 — ~21K source LOC across 116 files, 70 commits, 26 test files | Live: [all-things-agentic-allthingsage...](https://all-things-agentic-allthingsage.com) | Repo: github.com/Naman6019/all-things-agentic

Dual-Pipeline Multi-Agent System: built Careers and Studio workflows with Google ADK and LangGraph, ingesting ~2,900 postings/run with an 89.5% pre-LLM filter rate.

Guardrails & Full-Stack Delivery: shipped human send gates, Langfuse tracing, a 373-test hermetic suite (11s), and a Next.js dashboard on Google Cloud Run.

CareFlow Intelligence — Agentic RAG & Fine-Tuning for Clinical Research

Python, FastAPI, Next.js, Supabase/pgvector, ChromaDB, Unsloth, Ollama | ~7,300 LOC — synthetic data (MITRE Synthea) | Repo: github.com/Naman6019/CareFlow-Intelligence

RAFT Fine-Tuning & Local Serving: fine-tuned Qwen2.5-1.5B with reference evidence and hard distractors via 4-bit QLoRA (r=16, alpha=16), then served it locally through Ollama.

Agentic RAG & Ingestion: built bounded-tool citation-grounded RAG agent over 4 ingestion pipelines with Postgres GIN citation checks, plus a human-in-the-loop ingestion agent managing 18 FHIR R4 tables.

Human-in-the-Loop Data Agent: built a second agent for synthetic clinical data ingestion (10 patients, 479 encounters, 5,368 observations across 18 FHIR R4 tables) with mandatory human-approval gates; added an OpenRouter (Nemotron 3 Ultra) cloud fallback for the RAG agent.

Other Projects & Research

SQuAD QA Benchmarking & NLP Research: compared BERT, BiDAF, and DistilBERT for extractive QA; fine-tuned each with standardized eval pipeline (precision, recall, F1); published research paper in IJMDES. GitResume

Agent: self-built evidence-backed resume CLI (PyPI: git-resume-agent, ~4K LOC, 28 commits) that validates managed claims before compiling Word/PDF output.

EDUCATION & ACHIEVEMENTS

B.Tech, Computer Science Engineering (AI & ML Specialization) — VIT Amaravati

June 2021 – Sept 2025

CGPA: 8.2/10 | Published Research Paper in IJMDES | Certifications: Azure Administrator Associate, AWS Cloud Practitioner